

Stroke Prediction

1. Objective

Stroke is the 2nd leading cause of death globally, accounting for 11% of total deaths worldwide according to the World Health Organization (WHO).

The ability to predict the likelihood of a patient to get stroke based on the various health and patient demographics, like gender, age, various diseases, and smoking status, may facilitate preventative medical interventions to save lives.

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

2. Data

Data is from <https://www.kaggle.com/datasets/fedesoriano/stroke-prediction-dataset>

Data consists of

1) id: unique identifier 2) gender: "Male", "Female" or "Other" 3) age: age of the patient 4) hypertension: 0 if the patient doesn't have hypertension, 1 if the patient has hypertension 5) heart_disease: 0 if the patient doesn't have any heart diseases, 1 if the patient has a heart disease 6) ever_married: "No" or "Yes" 7) work_type: "children", "Govt_jov", "Never_worked", "Private" or "Self-employed" 8) Residence_type: "Rural" or "Urban" 9) avg_glucose_level: average glucose level in blood 10) bmi: body mass index 11) smoking_status: "formerly smoked", "never smoked", "smokes" or "Unknown" 12) *stroke*: 1 if the patient had a stroke or 0 if not Note: "Unknown" in smoking_status means that the information is unavailable for this patient

```
In [2]: df = pd.read_csv('healthcare-dataset-stroke-data.csv')
print(df.head())
df = df.set_index('id')
```

	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	Male	67.0	0	1	Yes	
1	51676	Female	61.0	0	0	Yes	
2	31112	Male	80.0	0	1	Yes	
3	60182	Female	49.0	0	0	Yes	
4	1665	Female	79.0	1	0	Yes	

	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	\
0	Private	Urban	228.69	36.6	formerly smoked	
1	Self-employed	Rural	202.21	NaN	never smoked	
2	Private	Rural	105.92	32.5	never smoked	
3	Private	Urban	171.23	34.4	smokes	
4	Self-employed	Rural	174.12	24.0	never smoked	

	stroke
0	1
1	1
2	1
3	1
4	1

```
In [3]: print('Number of patients: ', df.shape[0])
        print('Number of columns: ', df.shape[1])
```

```
Number of patients: 5110
Number of columns: 11
```

Cleaning data by removing duplicates.

```
In [4]: print('Duplicated observation ', df[df.duplicated()])
        df = df.drop_duplicates()
        print('New number of patients: ', df.shape[0])
```

```
Duplicated observation Empty DataFrame
Columns: [gender, age, hypertension, heart_disease, ever_married, work_type, Residence_type, avg_glucose_level, bmi, smoking_status, stroke]
Index: []
New number of patients: 5110
```

Information on data.

```
In [5]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 5110 entries, 9046 to 44679
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                 5110 non-null  object
1   age                    5110 non-null  float64
2   hypertension           5110 non-null  int64
3   heart_disease          5110 non-null  int64
4   ever_married           5110 non-null  object
5   work_type              5110 non-null  object
6   Residence_type         5110 non-null  object
7   avg_glucose_level      5110 non-null  float64
8   bmi                    4909 non-null  float64
9   smoking_status         5110 non-null  object
10  stroke                 5110 non-null  int64
dtypes: float64(3), int64(3), object(5)
memory usage: 479.1+ KB

```

In [6]: `df.describe()`

```

Out[6]:

```

	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
count	5110.000000	5110.000000	5110.000000	5110.000000	4909.000000	5110.000000
mean	43.226614	0.097456	0.054012	106.147677	28.893237	0.048728
std	22.612647	0.296607	0.226063	45.283560	7.854067	0.215320
min	0.080000	0.000000	0.000000	55.120000	10.300000	0.000000
25%	25.000000	0.000000	0.000000	77.245000	23.500000	0.000000
50%	45.000000	0.000000	0.000000	91.885000	28.100000	0.000000
75%	61.000000	0.000000	0.000000	114.090000	33.100000	0.000000
max	82.000000	1.000000	1.000000	271.740000	97.600000	1.000000

In [7]: `df.isnull().sum()`

```

Out[7]:
gender                0
age                   0
hypertension          0
heart_disease         0
ever_married          0
work_type             0
Residence_type        0
avg_glucose_level     0
bmi                   201
smoking_status        0
stroke                0
dtype: int64

```

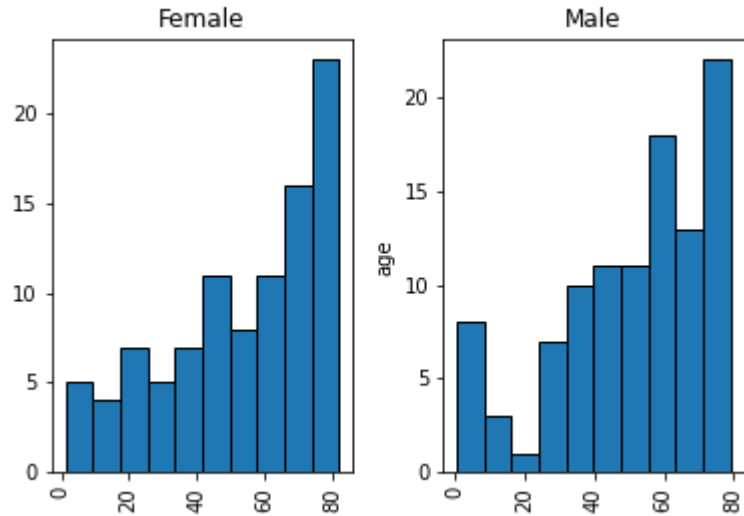
Feature Engineering: Imputation

BMI is body mass index, $\text{weight}/\text{height}^2$, and is a typical indicator of obesity and general health. As such, BMI could contribute strongly to risk of stroke.

Explore to identify any trends in age and gender for those null BMI to aid in imputation.

```
In [8]: df_bmi_null = df[df['bmi'].isnull() == True]
plt.figure();
df_bmi_null.hist(by='gender', column=['age'], grid=False, edgecolor = "black");
plt.ylabel('age');
```

<Figure size 432x288 with 0 Axes>



```
In [9]: print('Statistics of null BMI by Female gender')
print(df_bmi_null[df_bmi_null['gender'] == 'Female']['age'].describe())

print('Statistics of null BMI by Male gender')
print(df_bmi_null[df_bmi_null['gender'] == 'Male']['age'].describe())
```

Statistics of null BMI by Female gender

count	97.000000
mean	53.302268
std	22.519195
min	1.320000
25%	38.000000
50%	60.000000
75%	73.000000
max	82.000000

Name: age, dtype: float64

Statistics of null BMI by Male gender

count	104.000000
mean	50.880385
std	22.091668
min	0.480000
25%	35.000000
50%	57.500000
75%	70.250000
max	79.000000

Name: age, dtype: float64

Wide range in age span for both genders with null BMI values.

Given BMI varies substantially by age, need to calculate the median BMI for each age group for imputation of BMI. Add new column to classify each patient by their age group and impute the null BMI with their respective medians.

Age groups:

- less than 14 years, children
- 15 to 24, youth
- 25 to 64, adults
- greater than 65, seniors

```
In [10]: # Function to define age group based on age, returns age group name
```

```
def define_age_group(age):  
    if age <= 14:  
        return 'child'  
    elif age > 14 and age <=24:  
        return 'youth'  
    elif age > 24 and age <=64:  
        return 'adult'  
    elif age > 64:  
        return 'senior'
```

```
In [11]: # apply function to add age_group definition
```

```
df['age_group'] = df.apply(lambda x: define_age_group(x['age']), axis=1)
```

```
In [12]: # calculate median bmi based on age group
```

```
median_bmi_age_group = df[['bmi', 'gender', 'age_group']].groupby(['age_group', 'gender']).median()
```

```
In [13]: # Function to impute median based on age group for null bmi
```

```
def impute_bmi(age_group_x, gender_x, bmi_x):  
    if pd.isnull(bmi_x):  
        return median_bmi_age_group.query('age_group == @age_group_x & gender == @gender_x')['bmi']  
    else:  
        return bmi_x
```

```
In [14]: # apply function to add age_group definition
```

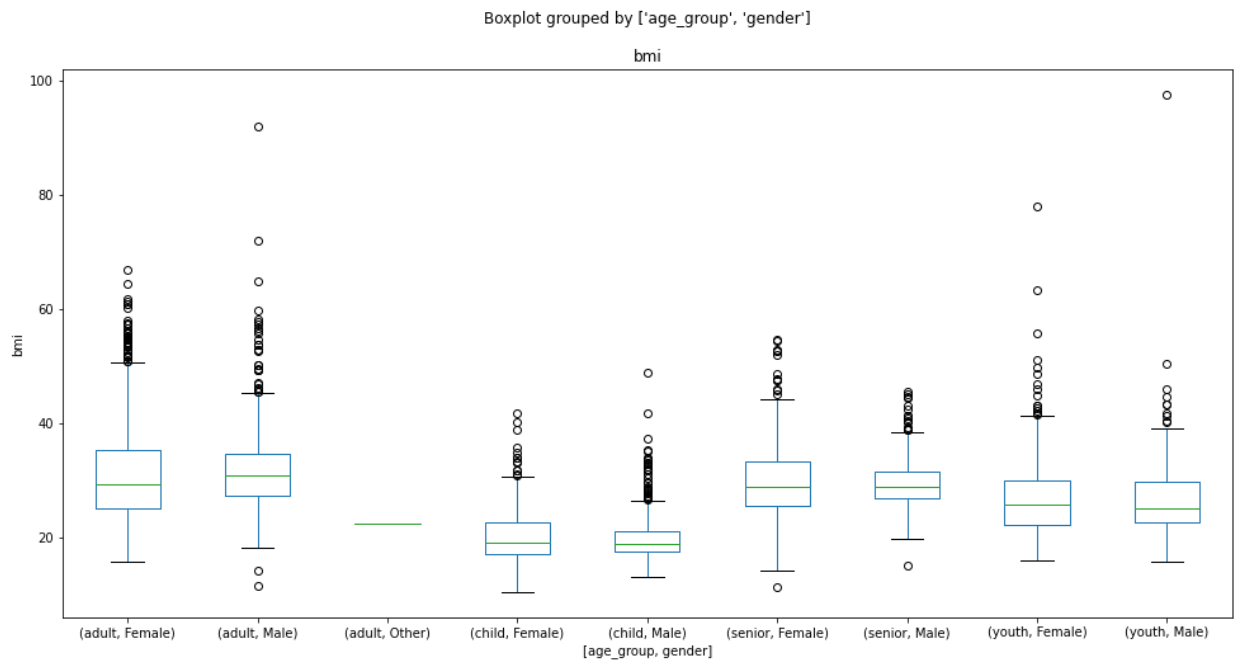
```
df['bmi'] = df.apply(lambda x: impute_bmi(x['age_group'], x['gender'], x['bmi']), axis=1)
```

```
# confirm no nulls in bmi  
df.isnull().sum()
```

```
Out[14]: gender          0  
age          0  
hypertension  0  
heart_disease 0  
ever_married  0  
work_type     0  
Residence_type 0  
avg_glucose_level 0  
bmi           0  
smoking_status 0  
stroke         0  
age_group      0  
dtype: int64
```

```
In [15]: # distribution of bmi by age group and gender
```

```
df.boxplot(by=['age_group', 'gender'], column=['bmi'], grid=False, figsize=(16,8));  
plt.ylabel('bmi');
```



Drop 'Other' gender as only 1 data point

```
In [16]: print(df['gender'].value_counts())
df = df[df['gender'] != 'Other']
print(df['gender'].value_counts())
```

```
Female    2994
Male      2115
Other         1
Name: gender, dtype: int64
Female    2994
Male      2115
Name: gender, dtype: int64
```

Summary statistics of data.

```
In [17]: df.describe()
```



```
Out[17]:
```




	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
count	5109.000000	5109.000000	5109.000000	5109.000000	5109.000000	5109.000000
mean	43.229986	0.097475	0.054022	106.140399	28.877471	0.048738
std	22.613575	0.296633	0.226084	45.285004	7.723424	0.215340
min	0.080000	0.000000	0.000000	55.120000	10.300000	0.000000
25%	25.000000	0.000000	0.000000	77.240000	23.700000	0.000000
50%	45.000000	0.000000	0.000000	91.880000	28.300000	0.000000
75%	61.000000	0.000000	0.000000	114.090000	32.800000	0.000000
max	82.000000	1.000000	1.000000	271.740000	97.600000	1.000000




Distribution of data.




Imbalanced target classes, more non-stroke than stroke.


```
In [18]: continuous_features = ['age', 'avg_glucose_level', 'bmi']
for column in df.columns:
    if column in continuous_features:
        plt.figure()
        df[column].plot.hist(edgecolor = "black");
        plt.xlabel(column)
    else:
        plt.figure()
        df[column].value_counts().plot(kind='bar')
        plt.xlabel(column)
```













```
In [19]: print('Number of non-stroke vs stroke')
df['stroke'].value_counts().plot(color=['green', 'red'])
print('Ratio of classes: ', df['stroke'].value_counts().iloc[1]/df['stroke'].value_cou
```

Number of non-stroke vs stroke
Ratio of classes: 0.05123456790123457




Visualize correlations between numerical data.

```
In [20]: sns.pairplot(df);
```



```
In [21]: plt.figure(figsize=(10,10))
sns.heatmap(df.corr(), annot=True)
```

```
Out[21]: <AxesSubplot:>
```



Rank strength of correlations

Highest correlation is only 0.33 between non-identical features, indicating not strong correlations between independent features individually.

```
In [22]: df.corr().unstack().sort_values(ascending=False)
```

```

Out[22]: age          age          1.000000
hypertension hypertension 1.000000
bmi          bmi          1.000000
avg_glucose_level avg_glucose_level 1.000000
heart_disease heart_disease 1.000000
stroke       stroke       1.000000
age          bmi          0.333143
bmi          age          0.333143
hypertension age          0.276367
age          hypertension 0.276367
heart_disease heart_disease 0.263777
heart_disease age          0.263777
stroke       age          0.245239
age          stroke       0.245239
avg_glucose_level avg_glucose_level 0.238323
avg_glucose_level age          0.238323
hypertension avg_glucose_level 0.174540
avg_glucose_level hypertension 0.174540
bmi          bmi          0.169657
bmi          avg_glucose_level 0.169657
hypertension bmi          0.162420
bmi          hypertension 0.162420
avg_glucose_level heart_disease 0.161907
heart_disease avg_glucose_level 0.161907
stroke       stroke       0.134905
stroke       heart_disease 0.134905
avg_glucose_level stroke       0.131991
stroke       avg_glucose_level 0.131991
hypertension stroke       0.127891
stroke       hypertension 0.127891
heart_disease hypertension 0.108292
hypertension heart_disease 0.108292
heart_disease bmi          0.041048
bmi          heart_disease 0.041048
stroke       stroke       0.039705
stroke       bmi          0.039705
dtype: float64

```

Feature engineering: categorical features

Convert categorical features (gender, ever_married, smoking_status, age_group, work_type, residence_type) to nominal.

```

In [23]: # collect list of categorical features
categorical_features = list()
for column in df.columns:
    if df[column].dtypes == 'object':
        categorical_features.append(column)

# remove age_group created
categorical_features.remove('age_group')

# convert categorical features to nominal
df_feat_eng = pd.get_dummies(df[categorical_features], drop_first=True)

# concat to existing numerical features
df_featured = pd.concat([df, df_feat_eng], axis=1)

```

```
# drop duplicate features
df_featured = df_featured.drop(['age_group', 'gender', 'ever_married', 'work_type', 'F
```

Scale feature values by min and max, similar to nominal features, so ranges are normalized to from 0 to 1.

```
In [24]: from sklearn.preprocessing import MinMaxScaler
mm = MinMaxScaler()
for column in df_featured.columns:
    df_featured[column] = mm.fit_transform(df_featured[[column]])

round(df_featured.describe().T, 3)
```

```
Out[24]:
```

	count	mean	std	min	25%	50%	75%	max
age	5109.0	0.527	0.276	0.0	0.304	0.548	0.744	1.0
hypertension	5109.0	0.097	0.297	0.0	0.000	0.000	0.000	1.0
heart_disease	5109.0	0.054	0.226	0.0	0.000	0.000	0.000	1.0
avg_glucose_level	5109.0	0.236	0.209	0.0	0.102	0.170	0.272	1.0
bmi	5109.0	0.213	0.088	0.0	0.153	0.206	0.258	1.0
stroke	5109.0	0.049	0.215	0.0	0.000	0.000	0.000	1.0
gender_Male	5109.0	0.414	0.493	0.0	0.000	0.000	1.000	1.0
ever_married_Yes	5109.0	0.656	0.475	0.0	0.000	1.000	1.000	1.0
work_type_Never_worked	5109.0	0.004	0.065	0.0	0.000	0.000	0.000	1.0
work_type_Private	5109.0	0.572	0.495	0.0	0.000	1.000	1.000	1.0
work_type_Self-employed	5109.0	0.160	0.367	0.0	0.000	0.000	0.000	1.0
work_type_children	5109.0	0.134	0.341	0.0	0.000	0.000	0.000	1.0
Residence_type_Urban	5109.0	0.508	0.500	0.0	0.000	1.000	1.000	1.0
smoking_status_formerly smoked	5109.0	0.173	0.378	0.0	0.000	0.000	0.000	1.0
smoking_status_never smoked	5109.0	0.370	0.483	0.0	0.000	0.000	1.000	1.0
smoking_status_smokes	5109.0	0.154	0.361	0.0	0.000	0.000	0.000	1.0

Split first into training and testing datasets

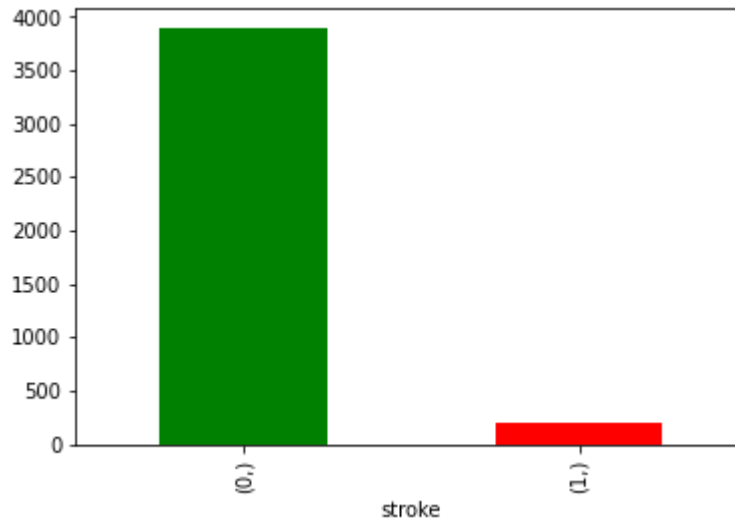
confirm class ratios for y_train and y_test with whole dataset.

```
In [25]: from sklearn.model_selection import train_test_split

X = df_featured.loc[:, df_featured.columns != 'stroke']
y = df_featured['stroke'].astype('int')
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, r

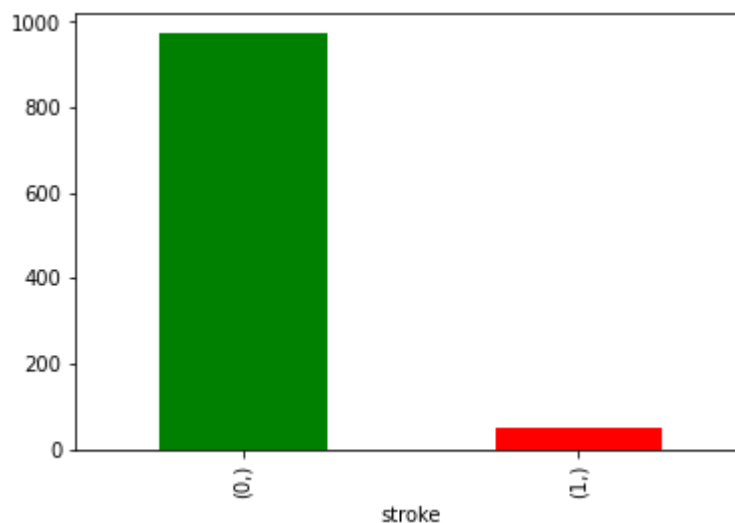
print('y_train')
print('Number of non-stroke vs stroke')
y_train.to_frame().value_counts().plot.bar(color=['green', 'red'])
print('Ratio of classes: ', y_train.to_frame().value_counts().iloc[1]/y_train.to_frame
```

```
y_train
Number of non-stroke vs stroke
Ratio of classes: 0.051183127572016464
```



```
In [26]: print('y_test')
print('Number of non-stroke vs stroke')
y_test.to_frame().value_counts().plot.bar(color=['green', 'red'])
print('Ratio of classes: ', y_test.to_frame().value_counts().iloc[1]/y_test.to_frame()
```

```
y_test
Number of non-stroke vs stroke
Ratio of classes: 0.051440329218107
```




Synthetic Minority Oversampling Technique (SMOTE)

SMOTE first creates many pairs or small clusters with two or more similar instances, the measure by instance distance such as Euclidean distance. Then, within the boundary of each pair or cluster, SMOTE uniformly permutes features value, one feature at a time, to populate a collection of similar synthesized instances within each pair or cluster. As a result, SMOTE creates a class-balanced synthetic dataset without adding duplicated instances with minority labels.

```
In [27]: from imblearn.over_sampling import RandomOverSampler, SMOTE
from imblearn.under_sampling import RandomUnderSampler
```




```
smote_sampler = SMOTE(random_state = 12345)
X_smo, y_smo = smote_sampler.fit_resample(X_train, y_train)
y_smo.value_counts().plot.bar(color=['green', 'red']);
```




Oversampling of the minority class to build up to the class size similar to that of majority

```
In [28]: oversample = RandomOverSampler(sampling_strategy='minority')
X_over, y_over = oversample.fit_resample(X_train, y_train)
y_over.value_counts().plot.bar(color=['green', 'red']);
```



Undersampling of majority class to reduce class size to similar to that of minority

```
In [29]: undersample = RandomUnderSampler(sampling_strategy='majority')
X_under, y_under = undersample.fit_resample(X_train, y_train)
y_under.value_counts().plot.bar(color=['green', 'red']);
```



3. Classifier models used to predict stroke,

1. Logistic Regression
2. SVC
3. Random Forest

- GridSearchCV is used for hyperparameter tuning of each of the 3 models.
- cross-validation of 5 folds to avoid overfitting.
- confusion matrix to visualize prediction/truth grid.

```
In [30]: import numpy as np
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_recall_fscore_support, confusion_matrix
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV, cross_val_score
```

Logistic Regression

confusion matrix and classification report to visualize and tabulation model performance.

```
In [31]: # fit to logistic regression model
param_grid = {'C':[0.001,0.1, 1,10, 100],
              'max_iter':[100, 500,1000]}

log = GridSearchCV(LogisticRegression(solver='lbfgs'),
                  param_grid,
                  cv=5)


log.fit(X_under,y_under)
print(log.best_params_)

# predict from X_test (which is imbalanced)
y_pred_log = log.predict(X_test)
confusion_log = confusion_matrix(y_test, y_pred_log)

plt.figure(figsize=(5, 5));
sns.heatmap(confusion_log, annot=True);
```

```
plt.xlabel("Predicted");
plt.ylabel("Actual");
```

```
{'C': 1, 'max_iter': 100}
```



```
In [32]: scores = list()

accuracy = accuracy_score(y_test, y_pred_log)
precision, recall, fbeta, support = precision_recall_fscore_support(y_test, y_pred_log)
auc = roc_auc_score(y_test, y_pred_log)
print(f"Accuracy is: {accuracy:.2f}")
print(f"Precision is: {precision:.2f}")
print(f"Recall is: {recall:.2f}")
print(f"Fscore is: {fbeta:.2f}")
print(f"AUC is: {auc:.2f}")

scores.append(('Logistic Regression', accuracy, precision, recall, fbeta, auc))
```

```
Accuracy is: 0.75
Precision is: 0.13
Recall is: 0.74
Fscore is: 0.63
AUC is: 0.75
```

Support Vector Machine Classifier

confusion matrix and classification report to visualize and tabulation model performance.

```
In [33]: # fit to logistic regression model
param_grid={'kernel':('linear', 'rbf'),
            'C':[0.01, 0.1, 1, 10, 100],
            'gamma':[0.00001, 0.0001, 0.01, 1]}
svc = GridSearchCV(SVC(probability=True),
                  param_grid,
                  cv=5)

svc.fit(X_under, y_under)
print(svc.best_params_)


# predict from X_test (which is imbalanced)
y_pred_svc = svc.predict(X_test)
```

```

confusion = confusion_matrix(y_test, y_pred_svc)
plt.figure(figsize=(5, 5));
sns.heatmap(confusion, annot=True);
plt.xlabel("Predicted");
plt.ylabel("Actual");

```

```
{'C': 10, 'gamma': 1e-05, 'kernel': 'linear'}
```



```

In [34]: accuracy = accuracy_score(y_test, y_pred_svc)
precision, recall, fbeta, support = precision_recall_fscore_support(y_test, y_pred_svc)
auc = roc_auc_score(y_test, y_pred_svc)
print(f"Accuracy is: {accuracy:.2f}")
print(f"Precision is: {precision:.2f}")
print(f"Recall is: {recall:.2f}")
print(f"Fscore is: {fbeta:.2f}")
print(f"AUC is: {auc:.2f}")

scores.append(('SVC', accuracy, precision, recall, fbeta, auc))

```

```

Accuracy is: 0.74
Precision is: 0.13
Recall is: 0.76
Fscore is: 0.64
AUC is: 0.75

```

Random Forest Classifier

confusion matrix and classification report to visualize and tabulation model performance.

```

In [35]: # fit to Random Forest model
param_grid = {'n_estimators': [50, 75, 100, 150, 200, 300]}
rf = GridSearchCV(RandomForestClassifier(random_state=12345),
                  param_grid,
                  cv=5)
rf.fit(X_under, y_under)
print(svc.best_params_)

# predict from X_test (which is imbalanced)
y_pred_rf = rf.predict(X_test)


```

```

confusion_rf = confusion_matrix(y_test, rf.predict(X_test))
plt.figure(figsize=(5, 5));
sns.heatmap(confusion_rf, annot=True);
plt.xlabel("Predicted");
plt.ylabel("Actual");

```

```
{'C': 10, 'gamma': 1e-05, 'kernel': 'linear'}
```



```

In [36]: accuracy = accuracy_score(y_test, y_pred_rf)
precision, recall, fbeta, support = precision_recall_fscore_support(y_test, y_pred_rf,
auc = roc_auc_score(y_test, y_pred_rf)
print(f"Accuracy is: {accuracy:.2f}")
print(f"Precision is: {precision:.2f}")
print(f"Recall is: {recall:.2f}")
print(f"Fscore is: {fbeta:.2f}")
print(f"AUC is: {auc:.2f}")

```

```
scores.append(('Random Forest', accuracy, precision, recall, fbeta, auc))
```

```

Accuracy is: 0.73
Precision is: 0.13
Recall is: 0.80
Fscore is: 0.67
AUC is: 0.76

```

Best model selection

All 3 models perform similarly with respect to accuracy, precision, recall, f-score and AUC for the test dataset. Since the test dataset is as imbalanced as the training dataset, the precision is quite low, which lowered the f-score. But the Random Forest model has slightly higher AUC and F-score, both of which are somewhat less sensitive to imbalanced dataset than precision and recall.

```

In [37]: df_scores = pd.DataFrame(scores, columns=['model', 'accuracy', 'precision', 'recall',
df_scores = df_scores.set_index('model')
df_scores

```

Out[37]:

	accuracy	precision	recall	fscore	AUC
model					
Logistic Regression	0.751468	0.133094	0.74	0.629581	0.746029
SVC	0.743640	0.131944	0.76	0.642393	0.751399
Random Forest	0.732877	0.132013	0.80	0.669672	0.764712

4. Key Findings

From the ranking of the feature importances of the Random Forest model below, the top 3 important features are age, avg_glucose_level and bmi. The rest of the features are much lower in terms of their importance.

This ranking is similar but not identical to the ranking for the linear correlations to stroke explored in the EDA above. For that list (see below), age, heart_disease and avg_glucose_level are the top 3 correlations, although they are low, at less than 0.25. BMI is the lowest at 0.0397.

The 2 evaluations indicate age, avg_glucose_level and bmi are significant features but they are not linearly correlated to stroke and a more complex relationship exists between these features and stroke prediction.

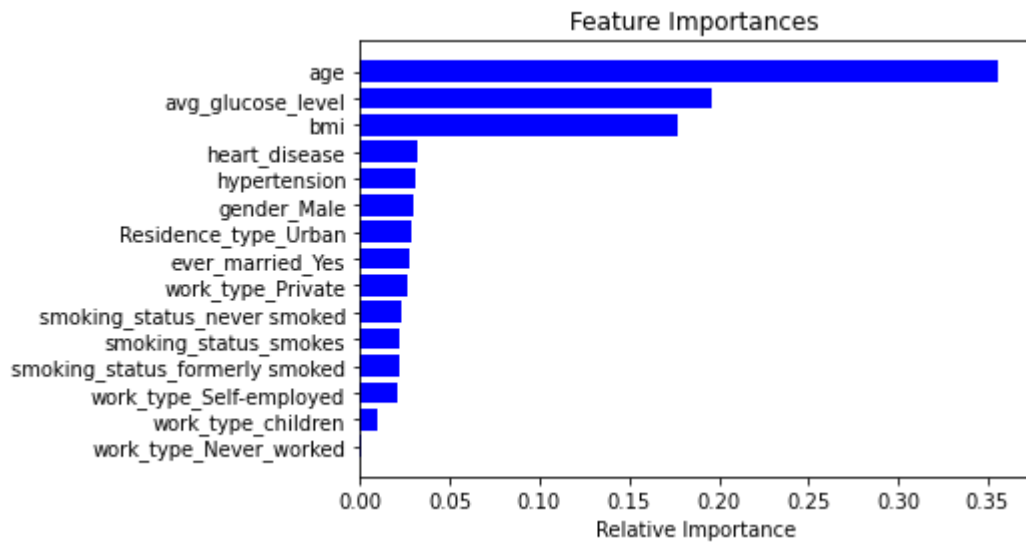
	correlation with stroke
age	0.245239
heart_disease	0.134905
avg_glucose_level	0.131991
hypertension	0.127891
bmi	0.039705

```
In [38]: # collect feature names
features = df_featured.columns.tolist()
features.remove('stroke')

# collect importances
importances = rf.best_estimator_.feature_importances_

# sort importances
indices = np.argsort(importances)

plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='b', align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```



5. Flaws and improvements

Stroke is a complex medical condition with a lot of other factors that may account for it beyond the ones explored in this dataset, namely age, gender, bmi, hypertension, smoking, etc.

From the exploratory analyses, age, bmi, hypertension and heart_disease have a slight correlation with stroke, but these are similar factors and appears to be co-related to themselves. But there are also other more apparent factors like occupation, pre-existing conditions, genetic predispositions, etc., which should be included to improve the prediction accuracies.

Further improvements can be with the use other classification models that uses boosting such as Gradient Boosted Classifier and AdaBoost.

In []: